

DevOps Interview Questions & Detailed Answers for Experienced Professionals

1. Explain the DevOps maturity model.

DevOps maturity measures how well an organization automates, collaborates, and continuously delivers software. Levels: Initial → Managed → Defined → Measured → Optimized. Higher maturity means more automation, observability, and rapid releases.

2. How do you design a CI/CD pipeline for microservices?

Use independent pipelines, automated testing, artifact versioning, containerization, deployment strategies like blue/green or canary, and centralized monitoring.

3. What is Infrastructure as Code and why is it important?

IaC automates environment provisioning using code (Terraform, CloudFormation). Benefits: consistency, version control, repeatability, and reduced manual errors.

4. Explain a real incident you handled in production.

Describe root cause analysis, logs/metrics analysis, communication with stakeholders, mitigation steps, and preventive measures like improved alerts or automation.

5. Blue-Green vs Canary deployment?

Blue-Green changes traffic instantly between two environments. Canary shifts traffic gradually to the new version, reducing risk.

6. How do you secure CI/CD pipelines?

Use secrets managers, role-based access control, no hard-coded credentials, artifact signing, vulnerability scanning, and audit logs.

7. Describe end-to-end Kubernetes architecture.

Includes API Server, Scheduler, Controller Manager, etcd, kubelet, kube-proxy, pods, deployments, services, ingresses, and storage classes.

8. How do you troubleshoot a failing Kubernetes pod?

Check pod logs, describe pod, validate image, environment variables, volume mounts, resource limits, and events. Use kubectl logs/describe.

9. How do you scale Kubernetes workloads?

Horizontal Pod Autoscaler, Vertical Pod Autoscaler, and Cluster Autoscaler. Scaling based on CPU, memory, or custom metrics.

10. Explain GitOps.

GitOps uses Git as a single source of truth for deployments. Tools like ArgoCD and Flux continuously sync cluster state with Git repos.

11. How do you handle secrets in DevOps?

Use AWS Secrets Manager, HashiCorp Vault, Kubernetes Secrets with encryption, or GitOps sealed secrets. Never store secrets in Git.

12. What is distributed tracing?

Tracing requests across microservices using tools like Jaeger or AWS X-Ray to identify latency and bottlenecks.

13. How do you design logging and monitoring?

Use centralized logging (ELK, Loki), metrics (Prometheus + Grafana), tracing, alerts, and dashboards for SLO/SLI monitoring.

14. How do you ensure high availability in cloud infrastructure?

Use multi-AZ deployment, load balancers, autoscaling, redundancy, managed databases, and distributed systems best practices.

15. What is Terraform state and why is it important?

Terraform state tracks infrastructure resources. Use remote state (S3 + DynamoDB) for team collaboration and locking.

16. How do you structure Terraform code for large environments?

Use workspaces, modules, DRY patterns, remote state, environment-specific variables, and versioned modules.

17. What is container orchestration?

Automated container management including deployment, scaling, load balancing, and healing. Kubernetes is the most popular orchestrator.

18. How do you optimize Docker images?

Use multi-stage builds, minimal base images, caching layers, .dockerignore, and avoid running containers as root.

19. Explain service mesh.

A service mesh (Istio, Linkerd) manages traffic, security, observability, and retries between microservices using sidecar proxies.

20. How do you handle zero-downtime deployments?

Use rolling updates, canary deployments, load balancing, readiness probes, and connection draining.

21. What is chaos engineering?

Testing system resilience using controlled failures (Chaos Monkey, Litmus). It improves fault tolerance and architecture stability.

22. What is SRE and how is it related to DevOps?

SRE applies engineering principles to operations. It focuses on SLOs, SLIs, error budgets, and reliability automation.

23. How do you containerize legacy applications?

Analyze dependencies, extract components, use minimal base images, create Dockerfile, and refactor if needed.

24. How do you secure Docker containers?

Use rootless containers, scan images, apply least privilege, restrict capabilities, and ensure image signing.

25. What strategies do you use for cost optimization in cloud?

Right-sizing, autoscaling, spot instances, lifecycle policies, serverless adoption, and shutting down unused resources.

26. How do you manage configuration drift?

Use IaC, continuous compliance checks, automation tools like Ansible, and immutable infrastructure.

27. What is multi-cloud deployment?

Running applications across multiple cloud providers for redundancy, cost optimization, or compliance.

28. How do you troubleshoot slow application performance?

Check CPU/memory, network latency, logs, metrics, tracing, database bottlenecks, and load balancer behavior.

29. What is artifact management?

Storing versioned build outputs in repositories like Nexus, Artifactory, or ECR for traceability and reusability.

30. How do you ensure disaster recovery?

Use backups, multi-region architecture, RPO/RTO planning, automated failover, and regular DR drills.

31. What is an ingress controller?

Manages external traffic into Kubernetes cluster using NGINX, Traefik, or AWS ALB.

32. How do you implement autoscaling?

Use metrics (CPU, memory, custom), HPA/VPA, and cloud autoscaling groups to dynamically adjust capacity.

33. What is a sidecar pattern?

Deploying helper containers alongside application containers for logging, proxying, or monitoring.

34. How do you enforce compliance in DevOps?

Use policy-as-code tools like OPA, Conftest, Terraform Sentinel, and continuous compliance scanners.

35. What is the CAP theorem?

Consistency, Availability, Partition Tolerance — only two can be fully achieved in distributed systems.

36. How do you test infrastructure code?

Use Terratest, Checkov, policy testing, and isolated test environments.

37. What is serverless DevOps?

Managing deployments for FaaS services (AWS Lambda) with CI/CD, monitoring, and versioning.

38. How do you secure Kubernetes clusters?

RBAC, pod security standards, network policies, runtime scanning, secrets encryption, and audit logs.

39. What is the difference between load balancer types?

Layer 4 LB handles TCP/UDP; Layer 7 LB handles HTTP/HTTPS with routing features.

40. How do you debug failed CI jobs?

Check logs, broken dependencies, environment mismatches, script errors, and failing tests.

41. How do you manage multiple environments?

Use separate IaC workspaces, namespaces, or clusters. Use Git branching strategies and automated promotions.

42. What is distributed logging?

Aggregating logs from multiple services using Fluentd, Logstash, or Loki.

43. How do you design a scalable infrastructure?

Use stateless apps, autoscaling, distributed caching, load balancers, and message queues.

44. What is the twelve-factor app?

A methodology for building scalable, cloud-native applications focusing on config, logs, build/release separation, and disposability.

45. How do you ensure environment consistency?

Use containers, IaC, automated provisioning, and centralized configuration.

46. What is horizontal vs vertical scaling?

Horizontal: adding more servers. Vertical: increasing CPU/RAM of existing ones.

47. What is container networking?

Overlay networks, CNI plugins, service discovery, and DNS resolution inside clusters.

48. How do you set up alerting?

Use thresholds, anomaly detection, SLO-based alerts, and integration with Slack or PagerDuty.

49. What is the role of a DevOps engineer in cloud architecture?

Automate provisioning, build CI/CD, ensure security, optimize performance, and manage releases.

50. Explain your biggest DevOps project.

Describe architecture, tools, automation improvements, CI/CD design, cloud setup, and final outcomes.